

STAT 435 HW 2

Pat McCornack

2022-10-14

1. Chapter 3: Question 14 of ISLR Book:

a) What is the form of the linear model y ?

The linear model y is of the form $y = \beta_1 x_1 + \beta_2 x_2 + \epsilon$ where $\beta_1 = 2$ and $\beta_2 = 0.3$ and the intercept is equal to 2.

```
library(MASS)
set.seed(1)
x1 <- runif(100)
x2 <- 0.5 * x1 + rnorm(100) / 10
y <- 2 + 2 * x1 + 0.3 * x2 + rnorm(100)
```


b) What is the correlation between x_1 and x_2 ?

The correlation coefficient between x_1 and x_2 is $r = 0.835$. This is strong positive correlation.

```
cor(x1, x2)
```

```
## [1] 0.8351212
```

```
plot(x1, x2)
```



c) Describe results of least squares regression.


The correlation coefficients for the linear model using x_1 and x_2 for least squares regression are: $\hat{\beta}_0 = 2.13$, $\hat{\beta}_1 = 1.44$, $\hat{\beta}_2 = 1.01$. The intercept is close, but both $\hat{\beta}_1$ and $\hat{\beta}_2$ are significantly different than the true coefficients of $\beta_1 = 2$ and $\beta_2 = 0.3$.


Since the p-value of $\hat{\beta}_1 = .049$ we can reject the null hypothesis $H_0 : \beta_1 = 0$ at a significance level $\alpha = 0.05$. We cannot reject the null hypothesis $\beta_2 = 0$ however.


```
lm.model <- lm(y ~ x1 + x2)
summary(lm.model)
```


```
##
## Call:
## lm(formula = y ~ x1 + x2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.8311 -0.7273 -0.0537  0.6338  2.3359
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.1305     0.2319   9.188 7.61e-15 ***
## x1              1.4396     0.7212   1.996  0.0487 *
## x2              1.0097     1.1337   0.891  0.3754
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##  
## Residual standard error: 1.056 on 97 degrees of freedom  
## Multiple R-squared:  0.2088, Adjusted R-squared:  0.1925  
## F-statistic: 12.8 on 2 and 97 DF,  p-value: 1.164e-05  
plot(lm.model)
```









d) Predict y using only x1

This model has a coefficient of $\hat{\beta}_1 = 1.98$ which is very close to the true value of 2. With a p-value of close to 0, we can reject the null hypothesis $\beta_1 = 0$.

```
x1_model <- lm(y ~ x1)
summary(x1_model)
```

```
##
## Call:
## lm(formula = y ~ x1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.89495 -0.66874 -0.07785  0.59221  2.45560
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.1124     0.2307   9.155 8.27e-15 ***
## x1              1.9759     0.3963   4.986 2.66e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.055 on 98 degrees of freedom
## Multiple R-squared:  0.2024, Adjusted R-squared:  0.1942
## F-statistic: 24.86 on 1 and 98 DF, p-value: 2.661e-06
```

e) Predict y using only x2

The coefficient $\hat{\beta}_2$ for this model is even further from the true value than the multivariate model at a value of 2.9 vs. the true value of 0.3, but the p-value **can** be used to reject the null hypothesis $\beta_2 = 0$ at a value of close to 0.

```
x2_model <- lm(y ~ x2)
summary(x2_model)

##
## Call:
## lm(formula = y ~ x2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.62687 -0.75156 -0.03598  0.72383  2.44890
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   2.3899      0.1949   12.26 < 2e-16 ***
## x2            2.8996      0.6330    4.58 1.37e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.072 on 98 degrees of freedom
## Multiple R-squared:  0.1763, Adjusted R-squared:  0.1679
## F-statistic: 20.98 on 1 and 98 DF,  p-value: 1.366e-05
```

f) Do these results contradict each other.

For x1: The coefficient $\hat{\beta}_1$ from (d) is closer to the true value than that from (c), but the results aren't necessarily contradictory.

For x2: The results for coefficient $\hat{\beta}_2$ from (c) and (e) are contradictory. In (c) the result for $\hat{\beta}_2$ did not allow us to reject the null hypothesis as it was distant from the true value of β_2 . The result for $\hat{\beta}_2$ from (e) did allow us to reject the null hypothesis despite being more distant from the true value.

g) Refit the model with an additional (mismeasured) observation

For the multivariate model the new observation made the coefficient less close to the true values and flipped the significance of the variables using the p-values. This model has that the null hypothesis can be rejected using the second coefficient but not the first.


In this model the observation is neither an outlier (defined as having a studentized residual $> \text{abs}(3)$) nor a high leverage point (based on Cook's Distance).


```
x1 <- c(x1, 0.1)
x2 <- c(x2, 0.8)
y <- c(y, 6)


lm.model <- lm(y ~ x1 + x2)
summary(lm.model)


##
## Call:
## lm(formula = y ~ x1 + x2)
##
```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.73348 -0.69318 -0.05263  0.66385  2.30619
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   2.2267     0.2314   9.624 7.91e-16 ***
## x1             0.5394     0.5922    0.911  0.36458
## x2             2.5146     0.8977    2.801  0.00614 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.075 on 98 degrees of freedom
## Multiple R-squared:  0.2188, Adjusted R-squared:  0.2029
## F-statistic: 13.72 on 2 and 98 DF,  p-value: 5.564e-06
plot(lm.model)
```










```
# Check for outliers  
studentized_resi = studres(lm.model)  
plot(studentized_resi)
```



```
# Check for leverage points
cd = cooks.distance(lm.model)
leverage_indices = which(cd > 4/nrow(lm.model))
leverage_indices
```

```
## integer(0)
```


In the model using only x1 the new observation made the coefficient more distant from the true value and raised the p-value, although the p-value is still significant enough to reject the null hypothesis.


The new observation is neither an outlier nor a high leverage point in this model either based on studentized residuals and cooks distance.


```
x1_model <- lm(y ~ x1)
summary(x1_model)
```


```
##
## Call:
## lm(formula = y ~ x1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.8897 -0.6556 -0.0909  0.5682  3.5665
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   2.2569     0.2390   9.445 1.78e-15 ***
## x1             1.7657     0.4124   4.282 4.29e-05 ***
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.111 on 99 degrees of freedom
## Multiple R-squared:  0.1562, Adjusted R-squared:  0.1477
## F-statistic: 18.33 on 1 and 99 DF,  p-value: 4.295e-05
plot(x1_model)
```










```
# Check for outliers
studentized_resi = studres(x1_model)
plot(studentized_resi)
```

```
outliers = which(studentized_resi > abs(3))
outliers
```

```
## 101
## 101
```

```
# Check for leverage points
cd = cooks.distance(x1_model)
leverage_indices = which(cd > 4/nrow(x1_model))
leverage_indices
```


```
## integer(0)
```


The results for the model based on only x_2 are the same. The coefficient was shifted further from its true value and the new observation is not an outlier or high leverage point based on studentized residuals and cooks distance.


```
x2_model <- lm(y ~ x2)
summary(x2_model)
```


```
##
## Call:
## lm(formula = y ~ x2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.64729 -0.71021 -0.06899  0.72699  2.38074
##
## Coefficients:
```

```
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.3451     0.1912  12.264 < 2e-16 ***
## x2          3.1190     0.6040   5.164 1.25e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.074 on 99 degrees of freedom
## Multiple R-squared:  0.2122, Adjusted R-squared:  0.2042
## F-statistic: 26.66 on 1 and 99 DF,  p-value: 1.253e-06
plot(x2_model)
```










```
# Check for outliers  
studentized_resi = studres(lm.model)  
plot(studentized_resi)
```



```
# Check for leverage points
cd = cooks.distance(lm.model)
leverage_indices = which(cd > 4/nrow(lm.model))
leverage_indices
```

```
## integer(0)
```

Q2

```
set.seed(1)
```

a)

Simulate the training data set.

```
n = 25

x <- rnorm(n,0,1)
eps <- rnorm(n,0,1)
y = exp(x) + eps
```

b)

Fit four regression models

```
p = 4

y1 <- lm(y ~ x)
y2 <- lm(y ~ x + x^2)
y3 <- lm(y ~ x + x^2 + x^3)
y4 <- lm(y ~ x + x^2 + x^3 + x^4)
```

c)

Create a training dataset with 500 observations

```
n = 500
x.test <- rnorm(n,0,1)
eps.test <- rnorm(n,0,1)
y.test <- exp(x.test) + eps.test
```

d)

Compute the test error for each of the four models.

```
MSE = c()

fitted.values <- coef(y1)[1] + x.test * coef(y1)[2]
MSE[1] <- mean((y.test - fitted.values)^2)

fitted.values <- coef(y2)[1] + x.test * coef(y2)[2]
MSE[2] <- mean((y.test - fitted.values)^2)

fitted.values <- coef(y3)[1] + x.test * coef(y3)[2]
MSE[3] <- mean((y.test - fitted.values)^2)

fitted.values <- coef(y4)[1] + x.test * coef(y4)[2]
MSE[4] <- mean((y.test - fitted.values)^2)
```

e)

Which model is the ‘best fit’ model?

The model with the lowest MSE value and therefore best fit is the first model $y \sim x$. It’s surprising to me that the linear model is the best fit as I would have expected the polynomial function $x + x^2 + x^3$ to be the best fit.

```
which.min(MSE)
```

```
## [1] 1
```

Q3

Using the Hitters dataset from the ISLR package:

```
library(ISLR)
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-4
```

```

library(dplyr)

##
## Attaching package: 'dplyr'
## The following object is masked from 'package:MASS':
##
##     select
## The following objects are masked from 'package:stats':
##
##     filter, lag
## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v purrr   0.3.4
## v tibble  3.1.8      v stringr 1.4.1
## v tidyr   1.2.0      v forcats 0.5.2
## v readr   2.1.2
## -- Conflicts ----- tidyverse_conflicts() --
## x tidyr::expand() masks Matrix::expand()
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## x tidyr::pack()   masks Matrix::pack()
## x dplyr::select() masks MASS::select()
## x tidyr::unpack() masks Matrix::unpack()
data("Hitters")
attach(Hitters)

```

a.

Split the data into training/testing sets

```

df <- data.frame(filter(Hitters, !is.na(Hitters$Salary))) # Remove NAs
df <- df[,-c(14,15,20)] # remove factor variables for regression

train_n <- ceiling(nrow(df) * 3/4)
test_n <- nrow(df) - train_n

data.train <- df[1:train_n,]
data.test <- df[train_n+1:nrow(df),]

```

b.

Fit the model using least squares regression and report the test error.

The mean test error of the linear model is -45.84.

```

lm.model <- lm(data.train$Salary ~ ., data = data.train)
lm.fitted <- as.matrix(cbind(1, data.test[-17])) %*% coef(lm.model)

```



```
test_error <- data.test$Salary - lm.fitted
mean(test_error, na.rm = TRUE)
```

```
## [1] -45.83773
```

```
test_error
```

```
##           [,1]
## -Ron Hassey   -34.97217
## -Rickey Henderson  851.34713
## -Reggie Jackson -1277.27213
## -Ron Kittle    183.92963
## -Ray Knight    -99.66930
## -Rick Leach     14.03376
## -Rick Manning  155.65622
## -Rance Mulliniks  63.03653
## -Ron Oester    170.21510
## -Rey Quinones   17.21533
## -Rafael Ramirez  731.76379
## -Ronn Reynolds  131.61764
## -Ron Roenicke  -191.58751
## -Ryne Sandberg  -330.55230
## -Rafael Santana  35.64796
## -Rick Schu     -55.42647
## -Ruben Sierra   -63.52250
## -Roy Smalley    309.92911
## -Robby Thompson -384.36189
## -Rob Wilfong     73.43356
## -Robin Yount    -367.17417
## -Steve Balboni  -520.29508
## -Scott Bradley  -237.48535
## -Sid Bream      -687.11739
## -Steve Buechele  -51.50404
## -Shawon Dunston  -95.08424
## -Scott Fletcher -186.22611
## -Steve Garvey    -49.01759
## -Steve Jeltz     -232.87837
## -Steve Lombardozzi -326.22622
## -Spike Owen      30.49317
## -Steve Sax      -1000.94415
## -Tony Bernazard  -353.82798
## -Tom Brookens    -67.88535
## -Tom Brunansky   473.28443
## -Tony Fernandez  -390.50857
## -Tim Flannery    -231.07392
## -Tom Foley       -62.72669
## -Tony Gwynn      -63.72088
## -Terry Harper    204.29391
## -Tommy Herr      111.48830
## -Tim Hulett      260.48883
## -Terry Kennedy   516.41635
## -Tito Landrum    149.88776
## -Tim Laudner     -63.34817
## -Tom Paciorek    -251.14170
## -Tony Pena       329.91081
```

## -Terry Pendleton	-77.19860
## -Tony Phillips	-178.65899
## -Terry Puhl	609.59721
## -Ted Simmons	-407.12015
## -Tim Teufel	-55.13832
## -Tim Wallach	367.94431
## -Vince Coleman	31.62912
## -Von Hayes	49.81745
## -Vance Law	134.89628
## -Wally Backman	-63.20466
## -Wade Boggs	265.00286
## -Will Clark	-504.20045
## -Wally Joyner	-780.19729
## -Willie McGee	240.88200
## -Willie Randolph	-32.44173
## -Wayne Tolleson	82.78513
## -Willie Upshaw	-179.09485
## -Willie Wilson	376.70912
## NA	NA
## NA.1	NA
## NA.2	NA
## NA.3	NA
## NA.4	NA
## NA.5	NA
## NA.6	NA
## NA.7	NA
## NA.8	NA
## NA.9	NA
## NA.10	NA
## NA.11	NA
## NA.12	NA
## NA.13	NA
## NA.14	NA
## NA.15	NA
## NA.16	NA
## NA.17	NA
## NA.18	NA
## NA.19	NA
## NA.20	NA
## NA.21	NA
## NA.22	NA
## NA.23	NA
## NA.24	NA
## NA.25	NA
## NA.26	NA
## NA.27	NA
## NA.28	NA
## NA.29	NA
## NA.30	NA
## NA.31	NA
## NA.32	NA
## NA.33	NA
## NA.34	NA
## NA.35	NA

## NA.36	NA
## NA.37	NA
## NA.38	NA
## NA.39	NA
## NA.40	NA
## NA.41	NA
## NA.42	NA
## NA.43	NA
## NA.44	NA
## NA.45	NA
## NA.46	NA
## NA.47	NA
## NA.48	NA
## NA.49	NA
## NA.50	NA
## NA.51	NA
## NA.52	NA
## NA.53	NA
## NA.54	NA
## NA.55	NA
## NA.56	NA
## NA.57	NA
## NA.58	NA
## NA.59	NA
## NA.60	NA
## NA.61	NA
## NA.62	NA
## NA.63	NA
## NA.64	NA
## NA.65	NA
## NA.66	NA
## NA.67	NA
## NA.68	NA
## NA.69	NA
## NA.70	NA
## NA.71	NA
## NA.72	NA
## NA.73	NA
## NA.74	NA
## NA.75	NA
## NA.76	NA
## NA.77	NA
## NA.78	NA
## NA.79	NA
## NA.80	NA
## NA.81	NA
## NA.82	NA
## NA.83	NA
## NA.84	NA
## NA.85	NA
## NA.86	NA
## NA.87	NA
## NA.88	NA
## NA.89	NA

## NA.90	NA
## NA.91	NA
## NA.92	NA
## NA.93	NA
## NA.94	NA
## NA.95	NA
## NA.96	NA
## NA.97	NA
## NA.98	NA
## NA.99	NA
## NA.100	NA
## NA.101	NA
## NA.102	NA
## NA.103	NA
## NA.104	NA
## NA.105	NA
## NA.106	NA
## NA.107	NA
## NA.108	NA
## NA.109	NA
## NA.110	NA
## NA.111	NA
## NA.112	NA
## NA.113	NA
## NA.114	NA
## NA.115	NA
## NA.116	NA
## NA.117	NA
## NA.118	NA
## NA.119	NA
## NA.120	NA
## NA.121	NA
## NA.122	NA
## NA.123	NA
## NA.124	NA
## NA.125	NA
## NA.126	NA
## NA.127	NA
## NA.128	NA
## NA.129	NA
## NA.130	NA
## NA.131	NA
## NA.132	NA
## NA.133	NA
## NA.134	NA
## NA.135	NA
## NA.136	NA
## NA.137	NA
## NA.138	NA
## NA.139	NA
## NA.140	NA
## NA.141	NA
## NA.142	NA
## NA.143	NA

## NA.144	NA
## NA.145	NA
## NA.146	NA
## NA.147	NA
## NA.148	NA
## NA.149	NA
## NA.150	NA
## NA.151	NA
## NA.152	NA
## NA.153	NA
## NA.154	NA
## NA.155	NA
## NA.156	NA
## NA.157	NA
## NA.158	NA
## NA.159	NA
## NA.160	NA
## NA.161	NA
## NA.162	NA
## NA.163	NA
## NA.164	NA
## NA.165	NA
## NA.166	NA
## NA.167	NA
## NA.168	NA
## NA.169	NA
## NA.170	NA
## NA.171	NA
## NA.172	NA
## NA.173	NA
## NA.174	NA
## NA.175	NA
## NA.176	NA
## NA.177	NA
## NA.178	NA
## NA.179	NA
## NA.180	NA
## NA.181	NA
## NA.182	NA
## NA.183	NA
## NA.184	NA
## NA.185	NA
## NA.186	NA
## NA.187	NA
## NA.188	NA
## NA.189	NA
## NA.190	NA
## NA.191	NA
## NA.192	NA
## NA.193	NA
## NA.194	NA
## NA.195	NA
## NA.196	NA
## NA.197	NA

C.

Fit a ridge regression model on the training set, with λ chosen by cross-validation. Report test error obtained. The best lambda value ends up being .001 selected using cross validation. The mean test error of the ridge regression model is -46.2 and the vector of test errors are as follows:

```
# Cross-validation for lambda

lam <- seq(0.001, 2, length.out = 100)
k = 5
ncv = ceiling(nrow(data.train)/k)
cv.ind = rep(1:k, ncv)
cv.ind.rand = sample(cv.ind, nrow(data.train), replace = F)

cv.error <- c(); MSE.cv <- c()
for(i in 1:100){
  for(j in 1:k){
    l.train <- data.train[cv.ind.rand != j, ]
    y.train <- data.train$Salary
    ridge.model <- glmnet(data.train[-17], y.train, lambda = lam[i], alpha = 0)

    l.test <- data.train[cv.ind.rand == j, ]
    l.test.values <- l.test$Salary
    test.response <- as.matrix(cbind(1, l.test[-17])) %*% coef(ridge.model, s = lam[i])
    MSE.cv[j] = mean((l.test.values - test.response)^2)
  }
  cv.error[i] = mean(MSE.cv)
}
lam_index = which.min(cv.error)

ridge.model <- glmnet(data.train[-17], data.train$Salary, lambda = lam[lam_index], alpha = 0)

ridge.fitted <- as.matrix(cbind(1, data.test[-17])) %*% coef(ridge.model, s = lam[lam_index])
test_error <- data.test$Salary - as.numeric(ridge.fitted)
mean(test_error, na.rm = TRUE)

## [1] -46.19122

test_error

##      [1]   -36.79078   851.24480 -1278.96748   184.54370  -100.06800    13.41332
##      [7]   161.75241    60.24510   167.24070    16.03406   730.21029   130.90995
##     [13] -192.41663  -329.50197    31.78446   -55.79786   -63.33625   311.96959
##     [19] -387.19654    73.55683  -360.22149  -520.10163  -236.81900  -686.47443
##     [25]  -52.80200   -97.42345  -186.83451  -52.70070  -234.61467  -329.86623
##     [31]    30.62521 -999.40037  -352.31646  -63.52127   473.31713  -393.40336
##     [37] -232.17255  -63.14417  -67.41617   203.55431   114.75830   257.59755
##     [43]   517.53392   147.76785  -61.84191  -254.86225   328.94896  -77.68033
##     [49] -179.32041   608.56397  -408.76122  -56.24395   366.05661    30.61383
##     [55]    54.01019   134.83320  -65.19952   258.77512  -503.86038  -777.69172
##     [61]   241.47031  -28.11561    82.24028  -177.88565   378.76999         NA
##     [67]         NA         NA         NA         NA         NA         NA
##     [73]         NA         NA         NA         NA         NA         NA
##     [79]         NA         NA         NA         NA         NA         NA
##     [85]         NA         NA         NA         NA         NA         NA
##     [91]         NA         NA         NA         NA         NA         NA
```

## [97]	NA	NA	NA	NA	NA	NA
## [103]	NA	NA	NA	NA	NA	NA
## [109]	NA	NA	NA	NA	NA	NA
## [115]	NA	NA	NA	NA	NA	NA
## [121]	NA	NA	NA	NA	NA	NA
## [127]	NA	NA	NA	NA	NA	NA
## [133]	NA	NA	NA	NA	NA	NA
## [139]	NA	NA	NA	NA	NA	NA
## [145]	NA	NA	NA	NA	NA	NA
## [151]	NA	NA	NA	NA	NA	NA
## [157]	NA	NA	NA	NA	NA	NA
## [163]	NA	NA	NA	NA	NA	NA
## [169]	NA	NA	NA	NA	NA	NA
## [175]	NA	NA	NA	NA	NA	NA
## [181]	NA	NA	NA	NA	NA	NA
## [187]	NA	NA	NA	NA	NA	NA
## [193]	NA	NA	NA	NA	NA	NA
## [199]	NA	NA	NA	NA	NA	NA
## [205]	NA	NA	NA	NA	NA	NA
## [211]	NA	NA	NA	NA	NA	NA
## [217]	NA	NA	NA	NA	NA	NA
## [223]	NA	NA	NA	NA	NA	NA
## [229]	NA	NA	NA	NA	NA	NA
## [235]	NA	NA	NA	NA	NA	NA
## [241]	NA	NA	NA	NA	NA	NA
## [247]	NA	NA	NA	NA	NA	NA
## [253]	NA	NA	NA	NA	NA	NA
## [259]	NA	NA	NA	NA	NA	NA

d.

Fit the lasso model and report the test error along with the number of non-zero coefficients.

The mean error for the lasso regression model is -46.19. However, I believe there's a mistake in the cross-validation code for lambda as the lasso error should not be the same as the ridge error. I cannot figure out where the error is. All variables are non-zero with the chosen lambda, which is likely not to be the case.

```
# Cross-validation for lambda

lam <- seq(0.001, 2, length.out = 100)
k = 5
ncv = ceiling(nrow(data.train)/k)
cv.ind = rep(1:k, ncv)
cv.ind.rand = sample(cv.ind, nrow(data.train), replace = F)

cv.error <- c(); MSE.cv <- c()
for(i in 1:100){
  for(j in 1:k){
    l.train <- data.train[cv.ind.rand != j, ]
    y.train <- data.train$Salary
    lasso.model <- glmnet(data.train[-17], y.train, lambda = lam[i], alpha = 1)

    l.test <- data.train[cv.ind.rand == j, ]
    l.test.values <- l.test$Salary
    test.response <- as.matrix(cbind(1, l.test[-17])) %*% coef(lasso.model, s = lam[i])
```

```

    MSE.cv[j] = mean((l.test.values - test.response)^2)
  }
  cv.error[i] = mean(MSE.cv)
}
lam_index = which.min(cv.error) # which value of lambda

lasso.model <- glmnet(data.train[-17], data.train$Salary, lambda = lam[lam_index], alpha = 1)

lasso.fitted <- as.matrix(cbind(1, data.test[-17])) %*% coef(lasso.model, s = lam[lam_index])
test_error <- data.test$Salary - as.numeric(lasso.fitted)
mean(test_error, na.rm = TRUE)

```

```
## [1] -46.19173
```

```
test_error
```

```

## [1] -36.79955 851.23784 -1278.95716 184.54736 -100.07141 13.41706
## [7] 161.75804 60.24214 167.22721 16.02155 730.17474 130.90619
## [13] -192.40581 -329.46812 31.77112 -55.80791 -63.34646 311.95799
## [19] -387.20906 73.57332 -360.19466 -520.09559 -236.81069 -686.45748
## [25] -52.81484 -97.43921 -186.83005 -52.71564 -234.60490 -329.87339
## [31] 30.62946 -999.39629 -352.30814 -63.50388 473.30693 -393.41548
## [37] -232.16825 -63.14119 -67.45451 203.55331 114.80419 257.57301
## [43] 517.53739 147.76996 -61.83148 -254.86133 328.93473 -77.67423
## [49] -179.31176 608.55263 -408.78901 -56.24795 366.04420 30.59123
## [55] 54.03272 134.84826 -65.20677 258.74488 -503.87019 -777.67355
## [61] 241.46705 -28.08873 82.24204 -177.86735 378.78292 NA
## [67] NA NA NA NA NA NA
## [73] NA NA NA NA NA NA
## [79] NA NA NA NA NA NA
## [85] NA NA NA NA NA NA
## [91] NA NA NA NA NA NA
## [97] NA NA NA NA NA NA
## [103] NA NA NA NA NA NA
## [109] NA NA NA NA NA NA
## [115] NA NA NA NA NA NA
## [121] NA NA NA NA NA NA
## [127] NA NA NA NA NA NA
## [133] NA NA NA NA NA NA
## [139] NA NA NA NA NA NA
## [145] NA NA NA NA NA NA
## [151] NA NA NA NA NA NA
## [157] NA NA NA NA NA NA
## [163] NA NA NA NA NA NA
## [169] NA NA NA NA NA NA
## [175] NA NA NA NA NA NA
## [181] NA NA NA NA NA NA
## [187] NA NA NA NA NA NA
## [193] NA NA NA NA NA NA
## [199] NA NA NA NA NA NA
## [205] NA NA NA NA NA NA
## [211] NA NA NA NA NA NA
## [217] NA NA NA NA NA NA
## [223] NA NA NA NA NA NA
## [229] NA NA NA NA NA NA

```



```
## [235]      NA      NA      NA      NA      NA      NA
## [241]      NA      NA      NA      NA      NA      NA
## [247]      NA      NA      NA      NA      NA      NA
## [253]      NA      NA      NA      NA      NA      NA
## [259]      NA      NA      NA      NA      NA      NA
```

```
coef(lasso.model, s = lam[lam_index])
```

```
## 17 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              s1
## (Intercept) 120.0626842
## AtBat      -3.0760596
## Hits       10.9313223
## HmRun      -3.8202376
## Runs       -2.4366967
## RBI         0.7222072
## Walks       6.6604145
## Years       1.0011724
## CAtBat     -0.1547079
## CHits      -0.1705178
## CHmRun      0.9420960
## CRuns       1.6651054
## CRBI        0.8621874
## CWalks     -0.8744750
## PutOuts     0.4108652
## Assists     0.7001926
## Errors     -4.9637317
```

e

Comment on the obtained results.

The test error, with a mean of approximately -46, is fairly low given the scale of the numbers we're working with. The can pretty accurately predict salary using the various observations. The difference in test error between the linear and lasso/ridge regression was marginal.